A Methodological Framework for Emotionally Aware AI Companion Systems

Mahesh, **Sanjay**, **Harish**, **Sindhu**

Department of AI&DS, Parul University, India

2303031241200@paruluniversity.ac.in
2303031240034@paruluniversity.ac.in
2303031241357@paruluniversity.ac.in
2303031241244@paruluniversity.ac.in

Abstract

Emotional isolation is a common experience among individuals such as students living away from home and people lacking consistent interpersonal support. Recent advances in machine learning and artificial intelligence have facilitated the development of emotionally responsive systems; however, designing AI companions that are both behaviourally reliable and ethically grounded remains a significant challenge.

This paper introduces a methodological framework for an AI companion system that integrates visual perception, emotion recognition, and character-consistent interactions.The proposed framework ensures that responses remain emotionally aligned in real-time while explicitly avoiding the substitution of human relationships. Key design considerations include system architecture, training strategies, regularisation techniques, and parameter tuning approaches to achieve convergence stability and behavioural consistency.

Rather than focusing on large-scale deployment or quantitative performance metrics, this study emphasises methodological principles and design strategies for emotionally interactions.It serves as a foundational exploration of AI companion systems, laying the groundwork for future research involving full-scale implementation, empirical validation, and long-term interactions analysis.

Keywords: Emotional AI, AI Companions, Emotion Recognition, Human–AI Interactions, Ethical AI.

1. Introduction

Emotional isolation has emerged as a significant social and psychological challenge in modern society. Students living away from home, individuals working in highly digital environments, and those lacking consistent interpersonal support often experience reduced emotional engagement and limited opportunities for meaningful interactions. While digital communication platforms have enhanced connectivity, they do not necessarily provide emotional presence or empathy, resulting in a gap between communication frequency and emotional well-being. This gap has led to growing interest in artificial intelligence (AI) systems capable of emotionally Aware interactions.

Recent advances in machine learning, deep learning, and multimodal perception have enabled the creation of conversational agents that can process natural language, facial expressions, and contextual cues. Emotion recognition Models trained on textual, visual, and auditory data have demonstrated promising performance in identifying affective states. However, despite these technological advancements, many AI-driven conversational systems remain limited in emotional consistency, contextual continuity, and ethical grounding. While some systems can produce responses that appear emotionally appropriate in isolated exchanges, they often fail to maintain long-term emotional alignment or character coherence during extended interactions.

The development of AI companion systems, designed to support users through emotionally responsive dialogue, has gained traction in both academic research and commercial applications. Such systems can provide emotional support, reduce perceived loneliness, and facilitate reflective conversations. However, they also introduce notable challenges. Inconsistent emotional behaviour, abrupt shifts in tone, and over-personalisation can undermine user trust and emotional stability. Additionally, poorly designed systems risk encouraging emotional dependency or blurring the boundary between artificial interactions and genuine human relationships.

Training and optimising emotion-Aware AI models also present significant difficulties. These systems often require complex training strategies involving multimodal data, reinforcement learning, and memory-based architectures. Without proper regularisation and tuning, AI companions may exhibit unstable behaviour, exaggerated emotional responses, or contextually inappropriate outputs. Ensuring behavioural reliability and convergence stability is therefore a key challenge, particularly in real-time interactive scenarios.

Ethical considerations, however, further complicate the deployment of emotionally responsive AI. While such systems may offer temporary emotional support, they must avoid undermining human connections or manipulating users emotionally. Issues related to privacy, bias in emotion recognition, emotional manipulation, and long-term psychological effects demand careful attention. Consequently, emotionally Aware AI companions should be designed with clear functional boundaries, emotional moderation, and transparency.

In response to these challenges, this paper introduces a methodological framework for developing emotionally Aware AI companion systems that integrate visual perception, emotion recognition, and character-consistent interactions. The focus is on architectural design, training methodologies, and strategies that ensure emotional stability and ethical alignment, laying the foundation for future research and practical applications in real-time emotionally intelligent AI systems.

2. Related Work

Research on emotionally responsive artificial intelligence spans multiple domains, including conversational agents, emotion recognition, and human–AI interactions. Early conversational systems primarily relied on rule-based architectures, where predefined patterns were mapped to scripted responses. While such systems demonstrated basic conversational ability, they lacked contextual understanding, emotional Awareness, and adaptability, limiting their effectiveness in emotionally sensitive interactions.

With the advancement of machine learning and natural language processing, data-driven conversational agents emerged. Sequence-to-sequence models and transformer-based architectures enable more fluent and context-Aware dialogue generation. These systems improved linguistic coherence but, in practice, often treated emotional content as an implicit feature rather than an explicit design objective. As a result, emotional responses were frequently inconsistent, contextually shallow, or misaligned with user behaviour during prolonged interactions.

Emotion recognition has been studied extensively as a supporting component for emotionally Aware systems. Text-based emotion recognition approaches utilise linguistic features, sentiment analysis, and contextual embeddings to infer emotional states from user Input. While effective in controlled scenarios, purely text-based Models struggle with ambiguity, sarcasm, and emotionally neutral language. To address these limitations, multimodal emotion recognition approaches have been proposed, incorporating facial expressions, visual cues, and, in some cases, vocal features. Visual perception models using convolutional neural networks have shown potential in identifying affective states from facial expressions, yet they remain sensitive to lighting conditions, cultural variations, and dataset bias.

Several studies have explored the integration of emotion recognition into conversational agents. These systems typically adjust response tone or content, to some extent, based on detected emotional states. However, many existing implementations focus on short-term emotional appropriateness rather than long-term emotional alignment. Emotional state transitions across interactions are often handled implicitly or ignored, leading to abrupt shifts in system behaviour. This lack of emotional continuity reduces user trust and limits the perceived authenticity of AI companions.

AI companion systems represent a more specialised category of conversational agents designed for sustained interactions. Unlike task-oriented chatbots, AI companions aim to maintain engagement, empathy, and conversational continuity. Research in this area highlights challenges related to memory management, personality persistence, and emotional stability. Systems lacking explicit mechanisms for character consistency may generate contradictory responses over time, undermining emotional reliability. Consequently, character modelling and memory-Aware architectures have become areas of increasing interest.

Ethical considerations have also gained prominence in related work on emotionally responsive AI. Researchers have raised concerns regarding emotional dependency, privacy risks, and the psychological impact of prolonged interactions with AI companions. Some studies emphasise the importance of transparency and emotional moderation to ensure that AI systems support users without encouraging substitution for human relationships. These concerns underline the need for design frameworks that explicitly incorporate ethical boundaries into system behaviour.

Training strategies for emotionally aware systems present additional challenges, particularly in large-scale or long-term interaction settings. Emotion recognition Models and dialogue generators often require large datasets and complex optimisation procedures. Without appropriate regularisation, such systems may overfit emotional patterns or amplify emotional expressions in undesirable ways. Recent work has explored curriculum learning, constrained reinforcement

learning, and behaviour regularisation as potential solutions; however, these techniques are often evaluated in isolation rather than as part of an integrated companion framework.

In summary, existing research demonstrates substantial progress in conversational AI and emotion recognition, yet several gaps remain. Many systems prioritise response-level emotional correctness while neglecting long-term emotional alignment, character consistency, and ethical safeguards. Furthermore, methodological discussions on training stability and behavioural reliability are often secondary to performance evaluation. These limitations motivate the need for a unified methodological framework that addresses emotional Awareness, system stability, and ethical design within a single cohesive architecture, as proposed in this work.

3. Problem Definition and Design Objectives

Despite significant progress in conversational artificial intelligence, the development of emotional Aware AI companion systems remains an open research challenge. Existing systems often demonstrate surface-level emotional responsiveness without maintaining consistent emotional behaviour over extended interactions.This limitation becomes particularly problematic in AI companion scenarios, where sustained engagement, emotional alignment, and user trust are critical.

3.1 Problem Definition

The primary problem addressed in this work is the lack of methodological frameworks that ensure emotional consistency, behavioural stability, and ethical grounding in AI companion systems. While emotion recognition Models can detect affective cues from text or visual Input, their integration into real-time conversational systems is frequently ad hoc. Emotional signals are often treated as transient features rather than persistent contextual states, resulting in fragmented emotional responses.

Another challenge lies in character inconsistency. Many conversational agents lack explicit mechanisms to maintain personality traits, conversational tone, or emotional boundaries across interactions.This can lead to contradictory behaviour, abrupt emotional shifts, or inappropriate emotional escalation. Such inconsistencies reduce user confidence and undermine the perceived reliability of AI companions.

Training instability further compounds these issues. Emotion-Aware systems typically rely on complex training pipelines involving multi-modal Inputs, sequential decision-making, and feedback-driven optimisation. Without appropriate regularisation and constraint mechanisms, these systems may converge to unstable behavioural patterns, including exaggerated empathy, emotional mirroring, or contextually inappropriate responses.

Ethical risks represent an additional dimension of the problem. Emotionally responsive systems, if improperly designed, may encourage emotional dependency or blur the distinction between artificial and human interactions. The absence of explicit ethical constraints during system design increases the risk of emotional overreach and misuse.

3.2 Design Objectives

To address the identified challenges, this work defines a set of design objectives for emotionally Aware AI companion systems:

1. **Emotional Alignment:**
   The system should generate responses that are aligned with the user's emotional state while maintaining appropriate emotional moderation. Emotional Awareness should support empathy without reinforcing negative emotional patterns.

2. **Behavioural Stability:**
   The system must exhibit consistent behaviour across extended interactions. Emotional responses should evolve smoothly over time rather than fluctuate abruptly between conversational turns.

3. **Character Consistency:**
   The AI companion should maintain a coherent conversational persona, including stable tone, values, and interactions style. Character consistency contributes to trust and long-term usability.

4. **Real-Time Interactions Capability:**
   Emotional perception and response generation should operate efficiently to support real-time interactions without noticeable latency.

5. **Ethical Safeguards:**
   The system must incorporate explicit boundaries to avoid emotional dependency and to reinforce the role of human relationships. Transparency and emotional restraint are treated as core design principles.

6. **Training Robustness:**
   Training strategies should promote stable convergence and prevent overfitting to emotional cues. Regularisation and tuning mechanisms are essential to ensure reliable behaviour under diverse interactions scenarios.

By formalising these objectives, the proposed framework establishes clear criteria for the design and evaluation of emotionally Aware AI companion systems. These objectives guide the architectural and methodological choices discussed in the following section, which presents the proposed system framework.

4. Proposed System Architecture

The proposed AI companion system is designed as a modular and extensible framework that integrates emotion perception, contextual understanding, and character-consistent response generation. The architecture emphasises emotional alignment, behavioural stability, and ethical boundaries while supporting real-time interactions. Rather than Relying on a monolithic model, the system adopts a layered design in which each component performs a distinct function and contributes to overall emotional reliability.

4.1 System Overview

At a high level, the system consists of four primary layers: Input perception, emotion interpretation, contextual memory and character modelling, and response generation. These layers operate sequentially while sharing contextual information to ensure coherence across interactions turns. Such modularisation allows individual components to be refined or replaced without destabilising the entire system.

User interactions begins with multimodal Input acquisition, which may include textual Input and visual cues. These Inputs are processed to infer emotional states and contextual intent. The resulting emotional representation is then combined with stored contextual and character information to guide response generation.

4.2 Input Perception Module

The Input perception module is responsible for collecting and preprocessing user interactions signals. Textual Input is analysed to extract semantic meaning and linguistic features relevant to emotional interpretation. When available, visual Input is processed to capture facial expressions and non-verbal cues that may indicate emotional state.

This module does not attempt to make final emotional decisions. Instead, it produces structured representations that preserve uncertainty and contextual ambiguity. By avoiding premature emotion classification, the system reduces the risk of misinterpretation and overly confident emotional responses.

4.3 Emotion Interpretation Module

The emotion interpretation module integrates signals from the perception layer to estimate the user's affective state. Emotional inference is treated as a probabilistic process rather than a deterministic classification task. This allows the system to represent mixed or uncertain emotional states, which are common in real-world interactions.

Rather than reacting solely to instantaneous emotional cues, the module maintains a short-term emotional context that captures recent interactions history. This temporal smoothing mechanism helps prevent abrupt emotional shifts and supports gradual emotional transitions across conversational turns.

4.4 Contextual Memory and Character Modelling

To ensure consistency over extended interactions, the system incorporates a contextual memory component and a character modelling layer. Contextual memory stores relevant information from previous interactions, including topics discussed, emotional trends, and interactions patterns. This memory is selectively updated to avoid excessive accumulation of sensitive or irrelevant data.

Character modelling defines the AI companion's interactions style, emotional boundaries, and response tendencies. By separating character traits from emotional inference, the system ensures that responses remain consistent even when user emotions fluctuate. This separation prevents emotional overfitting and reinforces a stable conversational persona.

4.5 Response Generation Module

The response generation module synthesises outputs based on emotional interpretation, contextual memory, and character constraints. Rather than optimising solely for emotional expressiveness, the system prioritises appropriateness, clarity, and moderation. Emotional alignment is achieved by adjusting tone, phrasing, and response structure rather than explicit emotional mimicry.

Ethical constraints are embedded within this module to prevent responses that encourage emotional dependency or substitute human relationships. For example, the system avoids exclusive language and reinforces user autonomy where appropriate.

4.6 Interactions Flow and Real-Time Considerations

The interactions flow is designed to support real-time responsiveness while maintaining emotional reliability. Lightweight representations and efficient inference pathways are prioritised to reduce latency. Additionally, fallback mechanisms are incorporated to handle uncertainty, such as defaulting to neutral or supportive responses when emotional signals are ambiguous.

By combining modular perception, emotion-Aware interpretation, memory-based context handling, and character-consistent generation, the proposed architecture provides a structured foundation for emotionally Aware AI companion systems. The design emphasises stability and ethical grounding, setting the stage for the training strategies and optimisation techniques discussed in the following section.

5. Training Strategies and Stability Considerations

The methodological choices adopted in this framework are influenced by practical machine learning principles that emphasise stable training, systematic error reduction, and controlled model behaviour. Rather than pursuing aggressive optimisation or performance maximisation, the training strategy prioritises reliability, interpretability, and gradual improvement. This philosophy aligns with widely accepted machine learning best practices that advocate structured experimentation and disciplined model tuning.

5.1 Emotion-Aware Training Philosophy

Emotionally Aware AI systems differ from conventional predictive Models in that their outputs directly influence user experience and emotional perception. Consequently, training objectives must balance emotional responsiveness with behavioural restraint. Instead of optimising solely for emotional expressiveness, the system is trained to minimise emotionally inappropriate responses while maintaining supportive interactions.

Emotional signals are treated as contextual guidance rather than absolute targets. This reduces sensitivity to noisy or ambiguous emotional cues and prevents exaggerated emotional reactions. By focusing on reducing harmful or unstable behaviours first, the system follows a conservative training approach that favours safety over expressiveness.

5.2 Bias–Variance Perspective in Emotional Modeling

From a bias–variance standpoint, emotionally Aware systems are prone to both underfitting and overfitting. Underfitting may result in emotionally neutral or detached responses, while overfitting can cause excessive empathy, emotional mirroring, or dependency-inducing behaviour. The proposed framework explicitly manages this trade-off by limiting emotional intensity and encouraging generalisable emotional patterns.

Model capacity is adjusted incrementally, and emotional response complexity is increased only when stable behaviour is observed. This gradual scaling strategy helps maintain control over emotional outputs and reduces unexpected behaviour during interactions.

5.3 Regularization for Behavioural Reliability

Regularization plays a critical role in ensuring behavioural stability. Rather than focusing exclusively on parameter-level regularization, the framework incorporates behaviour-level constraints. These constraints restrict extreme emotional responses, repetitive affirmations, or overly personalised language patterns.

By penalising unstable emotional shifts and reinforcing consistent response styles, regularization serves as a safeguard against erratic behaviour. This approach aligns with the principle of addressing the most significant sources of error first before pursuing finer optimisation.

5.4 Curriculum-Based Training Approach

Training is structured using a curriculum-based strategy, where the system is initially exposed to emotionally neutral or low-intensity interactions scenarios. As training progresses, emotionally complex and ambiguous cases are gradually introduced. This staged learning process allows the model to develop stable foundational behaviour before handling emotionally sensitive situations.

Such a progression reduces convergence instability and improves generalization across diverse interactions contexts. It also mirrors practical machine learning workflows that emphasise learning simpler patterns before addressing complex edge cases.

5.5 Tuning and Error Analysis

Rather than relying on extensive hyper parameter searches, tuning decisions are guided by targeted error analysis. Emotional misalignments, inconsistent tone, and contextually inappropriate responses are examined qualitatively to identify dominant failure modes. Adjustments are then applied to specific system components rather than globally modifying the entire model.

This targeted tuning approach reduces unintended side effects and supports incremental improvement. By systematically addressing one category of error at a time, the system achieves greater behavioural reliability without unnecessary complexity.

5.6 Stability over Performance Optimization

A key design decision in this framework is the prioritisation of stability over performance metrics. Since emotionally Aware AI companions are intended for sustained interactions, unpredictable behaviour poses a greater risk than reduced emotional expressiveness. Accordingly, training strategies favour predictable and ethically aligned behaviour even at the cost of reduced emotional intensity.

This conservative optimization philosophy ensures that the system remains dependable and trustworthy during real-time interactions, aligning with the broader goal of supporting users without replacing or dominating human relationships.

6. Ethical Considerations and System Limitations

The development of emotionally Aware AI companion systems raises important ethical considerations that must be addressed at the design level. Since such systems interact with users in emotionally sensitive contexts, improper behaviour can lead to unintended psychological, social, or ethical consequences. Accordingly, ethical safeguards are treated as a core component of the proposed framework rather than an afterthought.

6.1 Avoidance of Emotional Dependency

One of the primary ethical risks associated with AI companions is the potential for emotional dependency. Systems that exhibit excessive empathy, exclusivity, or emotional reassurance may unintentionally encourage users to rely on artificial interactions as a substitute for human relationships. To mitigate this risk, the proposed framework emphasises emotional moderation and avoids language that promotes exclusivity or long-term emotional reliance.

The system is designed to support users without positioning itself as a primary emotional authority. Where appropriate, responses reinforce user autonomy and encourage healthy real-world social interactions.

6.2 Transparency and Role Clarity

Maintaining transparency regarding the system's artificial nature is essential. Users should be Aware that emotional responses are generated algorithmically and do not represent genuine human emotion. Clear role boundaries help prevent emotional misinterpretation and reduce the likelihood of emotional manipulation.

The framework assumes explicit system identity and avoids anthropomorphic claims that could blur the distinction between artificial and human interactions.

6.3 Privacy and Data Sensitivity

Emotion recognition systems often process sensitive personal information, including emotional expressions and behavioural patterns. The proposed design minimises data retention by storing only essential contextual information and avoiding long-term storage of raw emotional data. Contextual memory is selectively updated and constrained to reduce privacy risks.

These considerations align with ethical principles of data minimisation and user consent, particularly in emotionally sensitive applications.

6.4 Bias and Misinterpretation Risks

Emotion recognition Models are susceptible to bias arising from training data limitations, cultural differences, and contextual ambiguity. Misinterpretation of emotional cues may lead to inappropriate responses or emotional misalignment. The framework addresses this risk by treating emotional inference probabilistically and incorporating uncertainty-Aware response strategies.

When emotional signals are ambiguous, the system defaults to neutral or supportive responses rather than making strong emotional assumptions.

6.5 System Limitations

Despite its structured design, the proposed framework has several limitations. The work focuses on methodological considerations rather than empirical validation, and the effectiveness of the framework has not been quantitatively evaluated. Additionally, real-world emotional complexity may exceed the representational capacity of current emotion recognition Models.

The framework also assumes cooperative user interactions and may not perform optimally under adversarial or deceptive Input scenarios. These limitations highlight the need for further empirical study and controlled deployment.

7. Discussion and Future Work

This work presents a conceptual and methodological exploration of emotionally Aware AI companion systems, emphasising stability, consistency, and ethical design. Rather than proposing a fully implemented system, the framework focuses on foundational design principles that address key shortcomings in existing emotionally responsive conversational agents. By prioritising emotional alignment over emotional intensity, the framework aims to support meaningful interactions while reducing behavioural unpredictability.

A central strength of the proposed approach lies in its modular architecture and conservative training philosophy. Separating emotion interpretation, character modelling, and response generation allows the system to maintain coherence even under uncertain emotional Input. The emphasis on stability-oriented training and targeted error analysis further contributes to predictable and reliable behaviour, which is critical in emotionally sensitive applications.

While the framework addresses several open challenges, it also opens multiple avenues for future research. Empirical validation using controlled user studies would provide insight into long-term emotional alignment and user trust. Future work may explore the integration of additional modalities, such as speech and physiological signals, to improve emotional inference accuracy. Longitudinal studies examining sustained interactions patterns could further inform the balance between emotional responsiveness and ethical restraint.

Additionally, future research may investigate adaptive character modelling strategies that evolve within predefined ethical boundaries. Exploring lightweight deployment strategies and on-device processing could also improve privacy preservation and real-time performance. These directions would contribute to the maturation of emotionally Aware AI companion systems from conceptual frameworks to responsibly deployed applications.

8. Conclusion

Emotionally Aware AI companion systems represent a promising yet challenging direction in human–AI interactions. While advances in machine learning have enabled emotionally responsive behaviour, ensuring long-term consistency, stability, and ethical alignment remains a critical concern. This paper has presented a methodological framework that integrates emotion perception, contextual memory, and character-consistent interactions to address these challenges.

By focusing on system design principles, training strategies, and ethical considerations, the proposed framework emphasises reliability over performance optimisation. The work deliberately avoids large-scale deployment claims and instead serves as a foundational study that highlights key design trade-offs in emotionally responsive AI systems.

The proposed framework contributes to ongoing discussions on responsible AI companionship by outlining practical methodological choices that support emotionally aligned interactions without replacing human relationships. As emotionally Aware systems continue to evolve, such design-oriented studies will play an essential role in guiding ethical and reliable development.

References

[1] R. W. Picard, *Affective Computing*. Cambridge, MA: MIT Press, 1997.

[2] S. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 4th ed. Pearson, 2021.

[3] D. Jurafsky and J. H. Martin, *Speech and Language Processing*, 3rd ed. Pearson, 2023.

[4] J. Devlin, M. W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," in *Proc. NAACL-HLT*, 2019.

[5] A. Vaswani et al., "Attention Is All You Need," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2017.

[6] F. Calefato, F. Lanubile, and N. Novielli, "Emotion Awareness in Software Engineering: A Systematic Review," *IEEE Trans. Affect. Comput.*, 2018.

[7] S. Poria, E. Cambria, R. Bajpai, and A. Hussain, "A Review of Affective Computing: From Unimodal Analysis to Multimodal Fusion," *Inf. Fusion*, 2017.

[8] P. Ekman, "An Argument for Basic Emotions," *Cognition and Emotion*, vol. 6, no. 3–4, pp. 169–200, 1992.

[9] S. D'Mello and J. Kory, "A Review and Meta-Analysis of Multimodal Affect Detection Systems," *ACM Comput. Surv.*, 2015.

[10] D. McDuff, R. El Kaliouby, T. Senechal, M. Amr, J. Cohn, and R. Picard, "Affectiva-MIT Facial Expression Dataset," *IEEE Trans. Affect. Comput.*, 2016.

[11] I. V. Serban et al., "A Survey of Available Corpora for Building Data-Driven Dialogue Systems," *Dialogue & Discourse*, 2015.

[12] E. Luger and A. Sellen, "Like Having a Really Bad PA: The Gulf between User Expectation and Experience of Conversational Agents," in *Proc. CHI Conf. Human Factors Comput. Syst.*, 2016.

[13] H. Y. Shum, X. He, and D. Li, "From ELIZA to XiaoIce: Challenges and Opportunities with Social Chatbots," *Front. Inf. Technol. Electron. Eng.*, 2018.

[14] E. M. Bender, T. Gebru, A. McMillan-Major, and S. Shmitchell, "On the Dangers of Stochastic Parrots," in *Proc. FAccT*, 2021.

[15] L. Floridi, J. Cowls, M. Beltrametti, et al., "AI4People—An Ethical Framework for a Good AI Society," *Minds Mach.*, 2018.

[16] B. D. Mittelstadt, P. Allo, M. Taddeo, S. Wachter, and L. Floridi, "The Ethics of Algorithms," *Big Data & Soc.*, 2016.

[17] S. Amershi et al., "Guidelines for Human–AI Interactions," in *Proc. CHI Conf. Human Factors Comput. Syst.*, 2019.

[18] V. Dignum, *Responsible Artificial Intelligence*. Cham: Springer, 2019.

[19] D. Silver et al., "Mastering the Game of Go with Deep Neural Networks and Tree Search," *Nature*, 2016.

[20] A. Ng, *Machine Learning Yearning*. Deeplearning.ai, 2018.